PROJECT 3 · AI GOVERNANCE

Model Intake Brief

A model ID in → an evidence brief out, in about 10 seconds

Sergii Rostovskyi · AI governance · Built with LangGraph

The problem, and who has it

WHO

AI governance lead or AI PMO

Approves third-party models before product teams build on them.

WHAT THEY DO TODAY

Reads the docs by hand

Model card, technical report, licence, usage policy and trust pages, one model at a time.

THE PAIN

37.4 h

a week spent assessing vendors, and 94% of companies cannot assess every vendor they want.

Why it matters now: since 2 August 2025 the EU AI Act requires providers of general-purpose AI models to give the companies that build on them documentation of what the model can and cannot do. Someone has to check that they did.

Sources: Whistic 2025 Impact Report, survey of 525 third-party risk leaders (all vendors, not AI only); Linklaters, "EU AI Act general-purpose AI model rules now apply", 2025.

How it works

One model ID in, one evidence brief out. No human steps in between.

1
Trigger
A model ID and one click (or one command)

→

2
Collect
Finds and reads the provider's public docs and model card

→

3
Check
EU AI Act and Code of Practice passages say what each of 18 checks needs

→

4
Write
The AI proposes a status and a quote per check; code verifies every quote and link

→

5
Deliver
A brief with gaps and provider questions, on screen and as a download

Collect and Check run at the same time. If a source fails, the brief says so instead of guessing.

HOW · LIVE DEMO

See it work

Qwen/Qwen3-8B → a verified brief with 18 checks

1 Start it

Live run on a real model ID

2 Watch it run

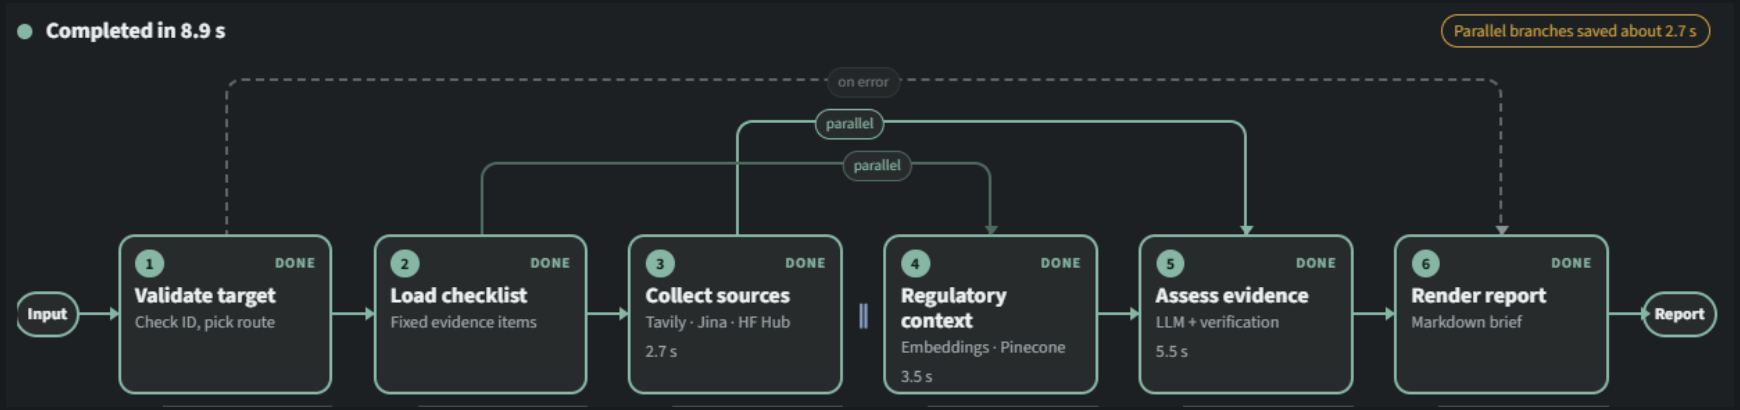
Data moves through the steps; two run in parallel

3 Open the brief

Coverage, gaps, provider questions

4 Prove it

Open one finding and click its source



Backup if the network or an API fails: the saved live reports and "Replay last run" work with no API calls.

What the user gets



Report overview from a real live run on Qwen/Qwen3-8B

~10 s

per brief, less than one cent of AI model cost

Sourced

every finding quotes the page and links to it; the quote is checked in the page text

Honest

shows what it could not find and turns each gap into a question for the provider

Under the hood

Part	What we use	Its job
Orchestration	LangGraph	Runs every step after the trigger, two in parallel
AI model	GPT-5.4 mini + text-embedding-3-small	Status and quote per check; embeds the rules
Data sources	Tavily · Jina Reader · Hugging Face Hub · link check	Finds, reads and re-checks the provider's docs
Grounding	Pinecone: EU AI Act + GPAI Code of Practice	Explains each check; never counts as evidence
Delivery	Web app + Markdown brief (and a command line)	Brief, dashboard and provider questions
Error handling	Error route, partial reports, quote and link checks	A tool fails or a quote is missing: no crash, no guessing

Why LangGraph: the core is verification logic, not app wiring. Code checks every quote and link, the graph state keeps each finding tied to its source, and it is all tested. n8n fits the delivery side later.

< $0.01
COST PER BRIEF*

9–18 s
TIME PER RUN

** AI model tokens only*

What’s next: go-to-market sprints

#	Goal and who	How we reach them	Deliverable	Success metric
1	**Prove it is useful** AI governance leads, AI PMO, risk reviewers	30-minute interviews + short survey	Revised demo + two annotated briefs	3 of 5 reviewers say it helps first-pass triage
2	**Pilot the handoff** AI inventory owners, governance ops	Guided handoff with one pilot team	Inventory-ready export map	80% of required fields mapped with no rework
3	**Test paid demand** Compliance leaders, procurement	Pricing talks + short pilot offer	Pricing hypothesis + pilot offer	2 prospects show purchase intent

The MVP is built: a working agent and two sample briefs. The sprints test whether reviewers value it, use it in their process and pay for it.



Thank you

Questions?

github.com/s-rostovskyi/model-intake-brief

Sergii Rostovskyi · Information, not legal advice